

activations

March 26, 2021

1 Activation Functions and their derivatives

[Reference blog](#)

{figure} ../images/artificial_neuron.gif :alt: Atifical Neuron Animation :class: bg-primary mb-1 :width: 400px Image Credit [<https://www.mql5.com/en/blogs/post/724245>]

1.1 Sigmoid Function

$$t = f(z) = \frac{1}{1 + e^{-z}}$$
$$\frac{dt}{dz} = f(z) * (1 - f(z)) = \frac{1 - e^{-z}}{1 + e^{-z}}$$


1.1.1 In Plain Python with Numpy

```
[1]: import matplotlib.pyplot as plt
import numpy as np

def sigmoid(x):
    s=1/(1+np.exp(-x))
    ds=s*(1-s)
    return s,ds
x=np.arange(-6,6,0.01)
sigmoid(x)
# Setup centered axes
fig, ax = plt.subplots(figsize=(9, 5))
ax.spines['left'].set_position('center')
ax.spines['right'].set_color('none')
ax.spines['top'].set_color('none')
ax.xaxis.set_ticks_position('bottom')
ax.yaxis.set_ticks_position('left')
# Create and show plot
ax.plot(x,sigmoid(x)[0], color="#307EC7", linewidth=3, label="sigmoid")
ax.plot(x,sigmoid(x)[1], color="#9621E2", linewidth=3, label="derivative")
ax.legend(loc="upper right", frameon=False)
fig.show()
```

<ipython-input-1-9103c13371b5>:21: UserWarning: Matplotlib is currently using module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot show the figure.

```
fig.show()
```




1.1.2 In Pytorch with torch.sigmoid()

“{warning}torch.range is deprecated and will be removed in a future release because its behavior is inconsistent with Python's range builtin. Instead, use torch.arange, which produces values in [start, end). Do not use: `x = torch.range(-6,6,0.01)`”

```
[2]: import torch
import numpy
import matplotlib.pyplot as plt

x = torch.arange(-6,6,0.01)
y = torch.sigmoid(x)

plt.plot(x.numpy(), y.numpy())
plt.show()
```




[AUTOGGRAD] (https://pytorch.org/tutorials/beginner/former_torchies/autograd_tutorial.html) is a core Torch package for automatic differentiation.

```
[6]: x = torch.arange(-6,6,0.01)
x.requires_grad_()
print(x[1:10])
t = torch.sigmoid(x)

t.backward(torch.ones(x.shape))

# x.grad the gradient at each x with respect to function t
dt = x.grad
plt.plot(x.detach().numpy(), dt.detach().numpy())
plt.show()
```

```
tensor([-5.9900, -5.9800, -5.9700, -5.9600, -5.9500, -5.9400, -5.9300, -5.9200,
        -5.9100], grad_fn=<SliceBackward>)
```



```
[27]: x
```

```
[27]: tensor([-6.0000, -5.9900, -5.9800, ...,  5.9700,  5.9800,  5.9900],
          requires_grad=True)
```

1.2 Hyperbolic tanh

$$t = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

$$\frac{dt}{dz} = 1 - t^2$$

1.2.1 In Plain Python and Numpy

```
[18]: import matplotlib.pyplot as plt
import numpy as np


def tanh(x):
    t=(np.exp(x)-np.exp(-x))/(np.exp(x)+np.exp(-x))
    dt=1-t**2
    return t,dt

z=np.arange(-4,4,0.01)
# print(tanh(z)[0])
tanh(z)[0].size,tanh(z)[1].size
```

```

# Setup centered axes
fig, ax = plt.subplots(figsize=(9, 5))
ax.spines['left'].set_position('center')
ax.spines['bottom'].set_position('center')
ax.spines['right'].set_color('none')
ax.spines['top'].set_color('none')
ax.xaxis.set_ticks_position('bottom')
ax.yaxis.set_ticks_position('left')
# Create and show plot
ax.plot(z,tanh(z)[0], color="#307EC7", linewidth=3, label="tanh")
ax.plot(z,tanh(z)[1], color="#9621E2", linewidth=3, label="derivative")
ax.legend(loc="upper right", frameon=False)
fig.show()

```



1.2.2 In Pytorch with torch.tanh()


```

[4]: import torch
import matplotlib.pyplot as pyplot

x = torch.arange(-6, 6, 0.01)
y = torch.tanh(x)

plt.plot(x.numpy(), y.numpy())
plt.show()


```



```
[5]: x = torch.arange(-6,6,0.01)
x.requires_grad_()
y = torch.tanh(x)

y.backward(torch.ones(x.shape))

plt.plot(x.detach().numpy(), x.grad.detach().numpy())
plt.show()
```



1.3 Relu

$$t = f(z) = \max(0, z)$$

$$\frac{dt}{dz} = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases} \quad (1)$$

““

```
[17]: import matplotlib.pyplot as plt
import numpy as np


def relu(x):
    t = [v if v >= 0 else 0 for v in x]
    dt = [1 if v >= 0 else 0 for v in x]
    t = np.array(t)
    dt = np.array(dt)
    return t,dt

z=np.arange(-4,4,0.01)
#print(relu(z)[0])
relu(z)[0].size,relu(z)[1].size
# Setup centered axes
fig, ax = plt.subplots(figsize=(9, 5))
ax.spines['left'].set_position('center')
ax.spines['bottom'].set_position('zero')
```

```

ax.spines['right'].set_color('none')
ax.spines['top'].set_color('none')
ax.xaxis.set_ticks_position('bottom')
ax.yaxis.set_ticks_position('left')
# Create and show plot
ax.plot(z,relu(z)[0], color="#307EC7", linewidth=3, label="tanh")
ax.plot(z,relu(z)[1], color="#9621E2", linewidth=3, label="derivative")
ax.legend(loc="upper right", frameon=False)
fig.show()

```




```

[8]: import torch
import matplotlib.pyplot as plt

relu = torch.nn.ReLU()
x = torch.arange(-6, 6, 0.01)
y = relu(x)

plt.plot(x.numpy(), y.numpy())
plt.show()

```

```
[15]: x = torch.arange(-6,6,0.01)
x.requires_grad_()
y = torch.nn.ReLU(x)

y.backward(torch.ones(x.shape))

plt.plot(x.detach().numpy(), x.grad.detach().numpy())
plt.show()
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-15-4743c799b6f4> in <module>
      3 y = torch.nn.ReLU(x)
      4
----> 5 y.backward(torch.ones(x.shape))
      6
      7 plt.plot(x.detach().numpy(), x.grad.detach().numpy())

C:\ProgramData\Anaconda3\envs\tf2\lib\site-packages\torch\nn\modules\module.py
  ~~~~ in __getattr__(self, name)
    945         if name in modules:
    946             return modules[name]
--> 947         raise AttributeError("'{}' object has no attribute '{}'.format
    948                               type(self).__name__, name))
    949
```

```
AttributeError: 'ReLU' object has no attribute 'backward'
```

1.4 Parametric and Leaky Relu

Dying ReLU The clipping effect of ReLU that helps with the vanishing gradient problem can also become an issue. Over time, certain outputs in the network can simply become zero and never revive again. This is called the ***dying ReLU*** problem. The Parametric ReLU (PReLU) is introduced to address the dying ReLU, where the leak coefficient a is a learned parameter.

$$t = \max(z, a * z) = f(a, z) \begin{cases} z & \text{if } z \geq 0 \\ a * z & \text{otherwise} \end{cases}$$
$$\frac{dt}{dz} = \begin{cases} 1 & \text{if } z \geq 0 \\ a & \text{otherwise} \end{cases}$$


```
[12]: import matplotlib.pyplot as plt
import numpy as np

def parametric_relu(a, x):
    t = [v if v >= 0 else a*v for v in x]
    dt = [1 if v >= 0 else a for v in x]
    t = np.array(t)
    dt = np.array(dt)
    return t, dt


z=np.arange(-6,6,0.01)
#print(relu(z)[0])
t,parametric_relu(0.05,z)
# Setup centered axes
fig, ax = plt.subplots(figsize=(9, 5))
ax.spines['left'].set_position('center')
ax.spines['bottom'].set_position('zero')
ax.spines['right'].set_color('none')
ax.spines['top'].set_color('none')
ax.xaxis.set_ticks_position('bottom')
ax.yaxis.set_ticks_position('left')
# Create and show plot
ax.plot(z,t[0], color="#307EC7", linewidth=3, label="tanh")
ax.plot(z,t[1], color="#9621E2", linewidth=3, label="derivative")
ax.legend(loc="upper right", frameon=False)
fig.show()
```

<ipython-input-12-42168a7efe00>:26: UserWarning: Matplotlib is currently using module://ipykernel.pylab.backend_inline, which is a non-GUI backend, so cannot

```
show the figure.  
fig.show()
```



```
[11]: import torch  
import matplotlib.pyplot as plt  
  
prelu = torch.nn.PReLU(num_parameters=1)  
x = torch.arange(-6, 6, 0.01)  
y = prelu(x)  
  
plt.plot(x.numpy(), y.detach().numpy())  
plt.show()
```



[]:

1.5 Softmax

The softmax function squashes the output of each unit to be between 0 and 1, like sigmoid function. However, the softmax operation also divides each output by the sum of all outputs, which gives a discrete probability distribution over k possible classes.

$$\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^k e^{x_j}}$$

```
[14]: import torch.nn as nn
import torch

softmax = nn.Softmax(dim=1)
x_input = torch.randn(1, 3)
y_output = softmax(x_input)
print(x_input)
print(y_output)
print(torch.sum(y_output, dim=1))
```

```
tensor([[1.4162, 0.9940, 1.4615]])
tensor([[0.3701, 0.2426, 0.3872]])
tensor([1.])
```

For more, [see this blog on Activation Functions in Neural Networks](#).